

Defining Schemas for the Grid Information Services

A proposal to be discussed as part of the Gridforum

Gregor von Laszewski and Peter Lane

Mathematics and Computer Science Division at Argonne National Laboratory
9700 S. Cass Avenue
Argonne, IL 60439 U.S.A.

Steve Fitzgerald

University of Southern Ccalifornia Information Sciences Institute
4676 Admiralty Way, Suite 1001
Marina del Rey, CA 90292-6695
U.S.A.

Brett Didier and Karen Schuchardt

Battelle, Pacific Northwest National Laboratory
Richland, WA 99352
U.S.A.

DRAFT: This draft is a draft version, please send comments to gregor@mcs.anl.gov
Editor: Gregor von Laszewski

Contents

1	Scope	2
2	Introductory Example	2
2.1	Objects	2
2.2	Aggregates	3
3	Grammar	5
3.1	Grid Object Specification	5
3.2	Class	5
3.3	Object Class	5
3.4	Object and Aggregate Names	6
3.5	Aggregate Class	6
3.6	Namespace	6
3.7	Attributes	6
3.8	Basic tokens	7
3.9	Comments	7
3.10	Definitions and their relevance to LDAP	7
4	Citation of attributes	8
5	Conversion Tools	8
5.1	gosConverter	9
A	ℒ_T_EX Binding	11
B	ASCII Graph Notation	12

1 Scope

The Grid Object Specification (GOS) format defines formal syntax for the definition of objects that form the core of the Grid Information Services (GIS). A formal syntax enables the definition of objects shared between a large number of participants. Furthermore, it allows the creation of automated tools for manipulating the objects. The GOS format includes some constructs that are not targeted for automation but documents policy issues, as well as, provides more complete information to the human reader.

29 Feb,2000

The format builds upon the syntax developed as part of the Globus/MDS project[1, 7] and has its roots in LDAP terminology[9]. A prototype implementation is developed as part of the CoG project [4].

The rest of this specification is structured as follows. First, we demonstrate the basic syntax on an example. Then we introduce gradually the detailed grammar, and finally we present a set of tools for manipulating the specifications.

2 Introductory Example

This example is intended to provide an intuitive overview of the GOS format while including all features of the format. In this example we define an object that represents a compute resource. Note, that the example in this document does not represent a complete and proper definition of the compute resource object. A complete proposal for such an object can be found elsewhere[6].

29 Feb,2000

2.1 Objects

```
Grid::ComputeResource OBJECT-CLASS ::= {
  SUBCLASS OF Grid::PhysicalResource
  RDN = hn (hostName)
  CHILD OF {
    Grid::organizationalUnit,
    Grid::organization,
  }
  MUST CONTAIN {
    canonicalSystemName :: cis,
    manufacturer :: cis, single,
    model :: cis, single,
    serialNumber :: ces , single,
  }
  MAY CONTAIN {
    diskDrive :: dn, multiple,
  }
}
```

Each object should have an appropriate description which explains its usage, the attributes, their types and the intended use. For example, we augment the attribute *serialNumber* with the following explanation:

Attributes

serialNumber: (ces) The serial number is used to distinguish the the compute resource form other products of the manufacturer/vendor.

...

Our object is associated with the namespace *Grid* and has the name *ComputeRe-*

29 Feb,2000

source. The object contains attributes that are grouped in a must contain, may contain sets. As the names indicate the must and may contain field define the validity of an object based on its attributes. Thus, if the attribute *serialNumber* is not defined in an object instantiation, the object would not be valid. Each attribute is identified by a name, a type, and a rule if the attribute can have multiple values (appear multiple times). In our case we see attributes for distinguished names, case sensitive, and case insensitive strings[8].

2.2 Aggregates

Next, we extend our example to include information related to geographical coordinates. Since such information can also be useful for other objects, it is useful to be able to define an aggregation of attributes that can be included in other objects. Thus, it is possible to keep the definition of the objects quite small in case repeated groups of attributes can be identified. Additionally, they are used to structure related attributes. We name such a specialized object class an *aggregate class*.

Note: Some Information Services do not allow to use multiple inheritance. In order to have an easy way of expressing this, aggregate objects are useful.

```
Grid::Location AGGREGATE-CLASS ::= {
  MUST CONTAIN {
    static :: cisboolean, single
  }
  MAY CONTAIN {
    latitude :: cisfloat, single
    longitude :: cisfloat, single
    locationName :: cis
  }
}
```

The description of the attributes could look like the following:

<i>static:</i>	(boolean) is true if the location of the object is fixed and can be associated with latitude and longitude values and false if the object is moving.
<i>locationName:</i>	(cis) e.g. state/city/.../floor/room/desk/...
<i>longitude:</i>	(cisfloat) relative east-west geographical position
<i>latitude:</i>	(cisfloat) relative north-south geographical position

Aggregate objects are not used by itself, but they are used to include all MUST and MAY CONTAIN attributes into an object in which they are referred. As shown next, we extend our example, the *ComputeResource* object, to include the aggregate object *Grid::location*.

```
Grid::ComputeResource OBJECT-CLASS ::= {
  SUBCLASS OF Grid::PhysicalResource
  RDN = hn (hostName)
  CHILD OF {
    Grid::organizationalUnit,
    Grid::organization,
  }
}
```

29 Feb,2000

Attributes

```

AGGREGATES {
    Grid::Location
}
MUST CONTAIN {
    canonicalSystemName :: cis,
    manufacturer :: cis, single,
    model :: cis, single,
    serialNumber :: ces , single,
}
MAY CONTAIN {
    diskDrive :: cis, multiple,
}
}

```

Besides the must and may contain groups we have also an AGGREGATES and a CHILD OF group. The aggregates group includes all objects which are “inlined” into the primary objects, that is all attributes of the aggregate objects are placed appropriately in the MAY or MUST CONTAIN groups. Note that there are a few restrictions which have to be kept in mind. First, the type, the multiplicity, and the presence (may or must) of attributes must be consistent in case an attribute is defined with the same name in multiple aggregate objects. In our case, the primary object Grid::ComputeResource would end up in an object with the following attributes:

```

Grid::ComputeResource OBJECT-CLASS ::= {
    SUBCLASS OF Grid::PhysicalResource
    RDN = hn (hostName)
    CHILD OF {
        Grid::organizationalUnit,
        Grid::organization,
    }
    MUST CONTAIN {
        canonicalSystemName :: cis,
        manufacturer :: cis, single,
        model :: cis, single,
        serialNumber :: ces , single,
        static :: cisboolean , single
    MAY CONTAIN {
        diskDrive :: cis, multiple,
        latitude :: cisfloat , single,
        longitude :: cisfloat , single,
        locationName :: cis
    }
}
}

```

In our case, the CHILD OF group indicates that the Grid::ComputeResource object can be stored logically in an Grid::organizationalUnit or and Grid::organization. The RDN or resources description name is a convenient abbreviation that is associated with the object. A primary object can be SUBCLASS of another primary object. In our case the Grid::ComputeResource object inherits all attributes that are defined by the CHILD OF, AGGREGATE, MUST CONTAIN, and MAY CONTAIN groups of the Grid::PhysicalResource.

Next, we will define the grammar more formally.

3 Grammar

The grammar is specified in Extended Backus Naur Form (EBNF). The following symbols are used to describe the grammar:

usage	semantics
(<i>a</i>)def	0 or 1 occurrences of <i>a</i>
(<i>a</i>) ⁺	1 or more occurrences of <i>a</i>
(<i>a</i>) [*]	0 or more occurrences of <i>a</i>
"T"	terminal
"#"	comments character

To simplify the grammar, multiple blanks between terminal symbols are not addressed. It is assumed that multiple blanks will be reduced to one.

3.1 Grid Object Specification

Each file following the GOS format contains one or more objects.

`<gos> ::= (<class>)+`

3.2 Class

A class can be either an object or an aggregate class.

`<class> ::= (<objectclass> | <aggregateclass>)`

3.3 Object Class

The rule for *object* is defined below. There are several constructs that comprise the object definition. It is important to note that an object in the schema may or may not be instantiated as a primary object. Primary objects will contain constructs such as RDN, CHILD OF and AGGREGATES where-as abstract and aggregate classes will not. Each construct is defined following the grammar specification.

03 Mar.,2000

```
<objectclass> ::= <objectname> "OBJECT-CLASS" " ::= " "{ "
    (SUBCLASS OF <objectname>)def
    (RDN = <attributename> "(" "
    <attributename> ")" )def
    (CHILD OF {
        (<objectname>;)+
    })def
    ( "AGGREGATES " "{ "
    (<aggregatename> " , " ){*}
    " } " )
    ( "MAY CONTAIN " "{ "
    (<attribute> " , " ){*}
    " } " )
    ( "MUST CONTAIN " "{ "
    (<attribute> " , " ){*}
    " } "
    " } " )
```

3.4 Object and Aggregate Names

The object and aggregate names are a simple concatenation of the namespace and the object name with the "::" character as delimitator.

03 Mar.,2000

```
<objectname> ::= <namespace> '::' ( <name> | <objectname> )
<aggregatename> ::= <name>
```

3.5 Aggregate Class

The rule for *aggregates* is defined below. An aggregate contains only the attributes defined under MAY and MUST CONTAIN. Aggregates do not use subclassing or inheritance.

03 Mar.,2000

```
<aggregateclass> ::= <objectname> "AGGREGATE-CLASS" "::=" "{ "
    ( "MAY CONTAIN " " { "
      (<attribute> " , " ) { * }
    " } " )
    ( "MUST CONTAIN " " { "
      (<attribute> " , " ) { * }
    " } "
    " } " )
```

3.6 Namespace

Namespaces are used to group a set of Object definitions together in order to avoid a name collision. Two special namespaces exist at the moment *X.500*, which indicates objects that are derived from the X.500 object specification, and *Grid*, which are objects that are building the core grid objects. Each application can define its own namespace. Nevertheless, we would like to point out that projects as part of the Grid activities may use the same abbreviation, thus we recommend to register the name with the Grid Information Services working group. A document on this issue will be prepared in future.

Aug.,1999

Namespaces are defined as part of the objectname. The rule for a namespace token is

```
<namespace> ::= <name>
```

Each namespace has a special object called ;namespace;::top. This object is inherited by all objects of this namespace.

3.7 Attributes

Attributes are defined in the following way:

```
<attribute> ::= <attributename> '::' <type>
<attributename> ::= <name>
```

In case an attribute is defined by an extended Schema the name is proceeded by the "ext" terminal.

Note:

```
<attributename> ::= <name> | 'ext' <name>
```

The complete list of LDAP types is defined in [2]. For our purposes we are mostly interested in the following types:

```
<type> ::= 'cis' | 'ces' | 'dn' | 'bin' | 'int' | <cis type>
```

We decided to include some additional types to make it easier for the user to decide upon their semantics, though these types are in reality of type *cis*. In future version of LDAP these types may be represented directly.

```
<cis type> ::= ``cisfloat`` | ``cisdate`` | ``cisboolean``
```

- *cisdate* – This type represents date similar to the output of `date -u` except that the day of the week is not kept and the month is provided as an integer value (e.g. 4 20 20:00:34 GMT 1999)
- *cisfloat* – This type represents floating point numbers formatted as strings. This does NOT include exponential notation (e.g. 0.02, 12.5)
- *cisboolean* – this type represents boolean values (e.g. “true” or “false”).

3.8 Basic tokens

As basic token we define names, characters, and numbers. The “...” in the character and number definition carry the usual “range-of” connotation.

```
<name> ::= <character> (<character> | <number>){*}
<character> ::= "a" | ... | "Z"
<number> ::= "0" | ... | "9"
```

Each of the major constructs of the *object* grammar are now defined. Some of these are defined in literature. Others have been added to communicate policy information regarding the structure of the DIT or how objects are to be instantiated.

3.9 Comments

All characters starting from the terminal character “#” until the end of the line are considered to be part of a comment. A comment can belong to an object or an attribute. This is indicated by the position of the comment, that is, it follows the definition.

We distinguish

```
<comment> := "#" <asciicharacter>*
<attributecomment> := <comment>
<objectcomment> := <comment>
```

3.10 Definitions and their relevance to LDAP

It is useful to understand the relationship between the terms used in the grammar and terms used while describing the concepts of LDAP. Though the grammar introduced above is a general information model, it provides some nice features, that allow to map it easily into an information model based on LDAP terminology. As we will see in [5] we can map the same model into an XML description.

SUBCLASS OF Objects can be derived from each other to provide data inheritance.

This concept is referred to in LDAP as subclassing or object class inheritance. An object that subclasses another object inherits all of the required and optional attributes of the superclass. The special abstract class “top” is inherited by all objects. This class contains at least the required attribute “objectclass” to guarantee that each object contains at least one object class. Similarly we define the class “Grid:Top” which is inherited by most objects. It is important to note that LDAP does not support multiple inheritance, as defined in object oriented languages such as C++. Nevertheless, due to the support of aggregation an object can be a member of more than one instantiated object class.

CHILD OF The "CHILD OF" construct is used to define a policy regarding the possible location(s) of the object in the DIT. An object which is specified as a "CHILD OF" a class or classes can only be placed in the DIT under objects of the specified type. This construct represents policy information only and is not in any way enforced. The intent of this policy is to keep a well structured tree where objects can be readily found. This attribute is only defined for primary objects.

RDN As specified in [9], each component of a DN is called a Relative Distinguished Name (RDN). Since each RDN is unique it guarantees that each entry in the the LDAP directory has a unique name. The RDN field in the notation introduced in this paper, specifies only the last component of the distinguished name which is assembled by the position in the DIT. This attribute is only defined for primary objects.

AGGREGATES The "AGGREGATES" construct is used to define the complete set of objects that are created together to form the primary object. As mentioned previously, LDAP does not support multiple inheritance. However through the use of aggregation, the primary object will include the attributes of all the aggregate objects. Aggregate objects are not allowed to be structural objects, that is they cannot be instantiated directly as objects in the DIT. An example of an aggregate object is Grid::Location. It will act as an aggregate of several objects such as Grid::PhysicalResource when the latter is instantiated and thus supply its attributes to the primary object.

MAY CONTAIN Attributes which are listed under MAY CONTAIN are optional attributes. When an object is uploaded to the LDAP server and schema checking is enabled no error will be returned if the attribute is missing in an object instantiation.

MUST CONTAIN Attributes which are listed under MUST CONTAIN are required attributes. When an object is uploaded to the LDAP server and schema checking is enabled an error will be returned if the attribute is missing in an object instantiation.

4 Citation of attributes

A convenient way to cite attributes in other documents is to use a string based token, that is based on the namespace, the class, and the attribute. We use the following short form:

```
<namespace>::<objectclass name>:<attribute name>
```

Thus, the string

```
Grid::software:description
```

would refer to the attribute description, which is part of the object software in the namespace Grid.

5 Conversion Tools

The purpose of GOS is to provide a common basis to develop tools which rely on the existence of "Grid" objects. The relationship between the format specification, the output generated by a translation tool and the application using this output is illustrated

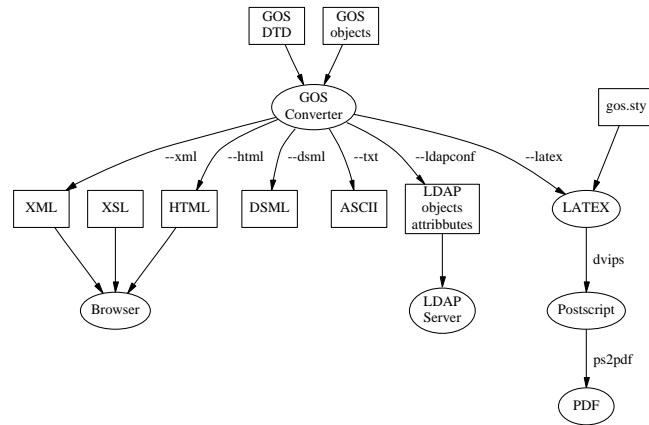


Figure 1: *The relationship between the good tools and the specification files*

in the figure below. An ASCII representation of this figure can be found in appendix B.

The advantage of separating the format specification from the actual implementation is to be able to utilize a variety of tools which allow the automatic generation of data structures and other useful but tedious tasks. A conversion tool should provide the following output formats:

- generate “at” and “oc” files, which are used at startup time of the LDAP server to initialize the schema. Often these two files are referred to in the slapd.conf file. See LDAP server documentation to find more out about how to set up an LDAP server.
- generate a set of HTML pages describing the GOS for convenient browsing of the specification. The URL to this page is defined in <http://www.globus.org/mds>.
- transform the LDAP based grammar into an XML based document, as well as, the reverse.

5.1 gosConverter

A prototype of a tool called gosConverter provides the following conversion options:

–input

The LDAP based specification of GOS is usually stored in a file called “grid-info-objectspecification.defs”.

–html

translates the input GOS schema definition to a set of HTML pages which can be viewed with a web browser.

–mdsml

translates the input GOS schema definition to mdsml a XML language that is an extension to DSML.

–dsml

translates the input GOS schema definition to DSML.

-ldapconfig

the GOS schema is translated into two files. These files store the attributes and the object classes which are used to initialize the slapd daemon.

Attributes File

The attributes are written into the file “grid-info-attributes.def.” This file is a slapd.ac.conf file. (to be consistent, we suggest renaming this file in future releases to “grid.slapd.at.conf”).

```
attribute o                cis
attribute ou               cis
attribute aci              ces
attribute objectName       cis
attribute ttl              cis
attribute userPassword     bin
attribute o                cis
attribute ou               cis
attribute email            cis
attribute phone            cis
attribute facsimiletelephonenumber cis
attribute gridadministrator cis
attribute seeAlso          dn
```

Object Class File

The attributes are written into the file “grid-info-objects.def.” This file is a slapd.oc.conf file. (To be consistent, we suggest renaming this file in future releases to “grid.slapd.oc.conf”).

```
objectclass GridTop
    requires
        objectclass,
        lastupdate
    allows
        aci,
        objectName,
        ttl

objectclass GridMDSAdministrator
    requires
        objectclass,
        lastupdate,
        userPassword,
        cn
    allows
        commonName,
        seeAlso,
        o,
        ou,
        aci,
        objectName,
        ttl
```

A L^AT_EX Binding

For convenience, it is possible to define the Grid Object Classes with the help of a L^AT_EX file. A conversion tool is available which can translate the L^AT_EX file into an XML object definition. Further manipulation tools are available which can be used to create from the XML objects a standard L^AT_EX description, a Web page, ASCII text, or a simple browser.

In order to utilize this convenient way of specifying the objects, please refer to the simple example given below.

```
%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%
\GRIDOBJECT{GridObjectExample}
%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%
```

The object {\em GridObjectExample} is used to

```
%-----
\begin{verbatim}
GridObjectExample OBJECT-CLASS ::= {
  SUBCLASS OF GridTop
  RDN = cn (commonName)
  CHILD OF {
    organization;
    organizationalUnit;
  }
  MUST CONTAIN {
    mustAttribute          :: cis ::singular,
  }
  MAY CONTAIN {
    mayAttributeA          :: cis,
    mayAttributeB          :: cis,
  }
}
\end{verbatim}

%-----
\begin{ATTRIBUTES}{AAAAAAAAA}

\item[mustAttribute] (cis,singular)
  This attribute stores ...

\item[mayAttributeA] (cis)
  This attribute stores ...

\item[mayAttributeB] (cis)
  This attribute stores ...

\end{ATTRIBUTES}
```

B ASCII Graph Notation

The tools graph in this paper is produced with dot[3].

```
digraph tools {

XML [shape=box];
XSL [shape=box];
HTML [shape=box];
DSML [shape=box];
ASCII [shape=box];
LDAP [shape=box,label="LDAP\nobjects\nattributes"];
GOSDTD [shape=box,label="GOS\nDTD"];
CONVERTER [label="GOS\nConverter"];
GOS [shape=box,label="GOS\nobjects"];
GOSSTY [shape=box,label="gos.sty"];
LDAPSERVER [label="LDAP\nServer"];

GOS      -> CONVERTER;
GOSDTD   -> CONVERTER;
CONVERTER -> XML      [label="--xml"];
CONVERTER -> DSML     [label="--dsml"];
CONVERTER -> HTML     [label="--html"];
CONVERTER -> ASCII    [label="--txt"];
CONVERTER -> LDAP     [label="--ldapconf"];
XML       -> Browser
XSL       -> Browser
HTML      -> Browser
LDAP      -> LDAPSERVER
CONVERTER -> LATEX     [label="--latex"];
GOSSTY    -> LATEX     ;
LATEX     -> Postscript [label="dvips"];
Postscript -> PDF       [label="ps2pdf"];
}
```

References

- [1] The Globus Object Class Definitions. http://www-isi.globus.org/mds/OCBrowser/globus_object_defs.html.
- [2] HOWES, T., KILLE, S., YEONG, W., AND ROBBINS, C. *The String Representation of Standard Attribute Syntaxes*, Mar. 1995. RFC 1778.
- [3] KOUTSOFIOS, E., AND NORTH, S. C. *Drawing Graphs with dot*, 1.5 ed. AT&T Labs - Research, Murray Hill, NJ, March 1999.
- [4] VON LASZEWSKI, G. The Globus CoG Web pages. <http://www.globus.org/cog>.
- [5] VON LASZEWSKI, G., , AND LANE, P. An XML binding for the Grid Object Specification. draft, Gridforum Information Services Working group, Feb. 2000.
- [6] VON LASZEWSKI, G., AND FITZGERALD, S. Representing Compute Resources. draft, Gridforum Information Services Working group, Feb. 2000.
- [7] VON LASZEWSKI, G., STELLING, P., FOSTER, I., KESSELMAN, C., AND C.LEE. A Fault Detection Service for Wide Area Distributed Computations. In *Proc. 7th IEEE Symp. on High Performance Distributed Computing* (July 1998), pp. pp. 268–278.
- [8] WAHL, M., COULBECK, A., HOWES, T., AND KILLE, S. *Lightweight Directory Access Protocol (v3): Attribute Syntax Definitions*, Dec. 1997. RFC 2252.
- [9] YEONG, W., HOWES, T., AND KILLE, S. *X.500 Lightweight Directory Access Protocol*, July 1993. RFC 1487.

Index

Attribute

- latitude, 3
- locationName, 3
- longitude, 3
- serialNumber, 2
- static, 3

gosConverter

- dsml, 9
- html, 9
- input, 9
- ldapconfig, 10
- mdsml, 9

- latitude, 3
- locationName, 3
- longitude, 3

- serialNumber, 2
- static, 3